

Deep Learning Assignment Report

1. Problem Understanding

- Task: Build a Convolutional Neural Network (CNN) for image classification using the CIFAR-10 dataset.

2. Model Design

- Architecture: 3 convolutional layers, 2 fully connected layers.
- Features: ReLU activation, Batch Normalization, and Dropout.

3. Results and Discussion

- Training/Validation Curves: Included.
- Test Accuracy: <Placeholder for test accuracy>.
- Confusion Matrix: Attached.

1. Activation Functions

(a) Sigmoid:

- **Definition:** A mathematical function that maps input values to a range between 0 and 1 using the formula:

$$f(x) = \frac{1}{1+e^{-x}}$$

- **Advantages:** Smooth gradient; interpretable output for probabilities.
- **Limitations:** Vanishing gradient problem; slow convergence; not zero-centered.
- **Use Case:** Logistic regression, binary classification problems.

(b) ReLU (Rectified Linear Unit):

Definition: Outputs the input directly if positive, otherwise output zero:

$$f(x) = \max(0, x)$$

- **Advantages:** Computationally efficient; mitigates vanishing gradient problems; promotes sparse activations.

- **Limitations:** Can suffer from the “dead neurons” issue.
- **Use Case:** Standard for most deep learning tasks, especially CNNs.

(c) Tanh:

- **Definition:** A hyperbolic tangent function that outputs values between -1 and 1:

$$f(x) = \frac{e^x - e^{-x}}{e^x + e^{-x}}$$

- **Advantages:** Zero-centered; better for centered data than Sigmoid.
- **Limitations:** Suffers from vanishing gradient issues.
- **Use Case:** Older neural network models; use has declined in favor of ReLU.

(d) Leaky ReLU:

- **Definition:** Allows small gradients for negative inputs by modifying the ReLU formula:

$$f(x) = \begin{cases} x, & x > 0 \\ ax, & x \leq 0 \end{cases}$$

- **Advantages:** Solves the dead neuron issue in ReLU.
- **Limitations:** Slightly more computationally expensive.
- **Use Case:** Alternative to ReLU in deeper networks.

2. Optimization Algorithms

Here's the explanation for your report:

(a) SGD (Stochastic Gradient Descent):

- **Description:** Updates weights using a single training example at each step.
Formula:

$$W = w - \eta \cdot \nabla L(w)$$

- **Advantages:** Simple; requires less memory.

- **Limitations:** Slow convergence; sensitive to learning rate.
- **Use Case:** Suitable for convex optimization problems.

(b) Adam (Adaptive Moment Estimation):

- **Description:** Combines the benefits of momentum and adaptive learning rates.
- **Advantages:** Fast convergence; adaptive to parameter scaling.
- **Limitations:** May not generalize well on some tasks.
- **Use Case:** Widely used in most modern deep learning tasks.

(c) RMSprop:

- **Description:** Maintains a running average of squared gradients to scale the learning rate.
- **Advantages:** Works well with non-stationary objectives.
- **Limitations:** Requires tuning of hyper parameters.
- **Use Case:** Commonly used for recurrent neural networks (RNNs)